

### Experiment 5

AIM: Write a program to perform weather forecasting using R

#### Code:

```
df<-read.csv('weather.csv')
apply(apply(df,2,is.na),2,sum)
df<-na.omit(df)
binca<-df$RainTomorrow
df$RainTomorrow[df$RainTomorrow == 'Yes']<-1
df$RainTomorrow[df$RainTomorrow == 'No']<-0
set.seed(123)
split = sample.split(df$RainTomorrow, SplitRatio = 0.8)
training_set = subset(df, split == TRUE)
test_set = subset(df, split == FALSE)
training_set<-training_set[-c(1,2,8,10,11,23)]
mytrainset2<-training_set[-18]
training_set<-lapply(training_set,as.numeric)
modell<-
glm(training_set$RainTomorrow~.,data=mytrainset2,family=binomial())
summary(modell)
test_set<-test_set[-c(1,2,8,10,11,23)]
mytestset2<-test_set[-18]
test_set<-lapply(test_set,as.numeric)
predi1<-predict(modell,mytestset2,type="response")
y_pred = ifelse(predi1 > 0.5, 1, 0)
tab1 <- table(Predicted = predi1, Actual = test_set$RainTomorrow)
missing_classerr <- mean(y_pred != test_set$RainTomorrow)
print(paste('Accuracy =', 1 - missing_classerr))
exp(coef(modell))
```

```

anova(model1,test="Chisq")

model2<-
glm(training_set$RainTomorrow~MinTemp+MaxTemp+Sunshine+WindGustS
peed+Humidity9am+Humidity3pm+Pressure3pm,data=mytrainset2,family=bin
omial())

summary(model2)

1-pchisq(deviance(model2)-deviance(model1),df.residual(model2)-
df.residual(model1))

predi2<-predict(model2,mytestset2,type="response")

binca<-ifelse(binca=="Yes",c(1),c(0))

summ=0

for(i in 1:length(predi1)){summ=summ+(binca[i]-predi1[i])^2}

summ

summ=0

for(i in 1:length(predi2)){summ=summ+(binca[i]-predi2[i])^2}

summ

mean(binca)

sum=0

for(i in 1:length(binca)){sum=sum+(binca[i]-mean(binca))^2}

sum

```

## Output:

```
> df<-read.csv('weather.csv')
> apply(apply(df,2,is.na),2,sum)
      Date      Location      MinTemp      MaxTemp      Rainfall      Evaporation      Sunshine
0          0          0          0          0          0          0          3
WindGustDir WindGustSpeed WindDir9am WindDir3pm WindSpeed9am WindSpeed3pm Humidity9am
3          2          31          1          7          0          0
Humidity3pm Pressure9am Pressure3pm Cloud9am Cloud3pm Temp9am Temp3pm
0          0          0          0          0          0          0
RainToday      RISK_MM RainTomorrow
0          0          0

> df<-na.omit(df)
> binca<-df$RainTomorrow
> df$RainTomorrow[df$RainTomorrow == 'Yes']<-1
> df$RainTomorrow[df$RainTomorrow == 'No']<-0
> set.seed(123)
> split = sample.split(df$RainTomorrow, SplitRatio = 0.8)
Error in sample.split(df$RainTomorrow, SplitRatio = 0.8) :
  could not find function "sample.split"
> training_set = subset(df, split == TRUE)
> test_set = subset(df, split == FALSE)
> training_set<-training_set[-c(1,2,8,10,11,23)]
> mytrainset2<-training_set[-18]
> training_set<-lapply(training_set,as.numeric)
Warning message:
In lapply(training_set, as.numeric) : NAs introduced by coercion
> model1<-glm(training_set$RainTomorrow~.,data=mytrainset2,family=binomial())
> summary(model1)

Call:
glm(formula = training_set$RainTomorrow ~ ., family = binomial(),
    data = mytrainset2)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-2.14927  -0.38430  -0.18007  -0.06652   2.43614

Coefficients:
(Intercept) 123.625886  54.710756  2.260 0.023844 *
MinTemp     -0.207058  0.158639 -1.305 0.191820
MaxTemp     -0.028772  0.266982 -0.108 0.914181
Rainfall    -0.037974  0.082738 -0.459 0.646255
Evaporation  0.149512  0.139944  1.068 0.285351
Sunshine    -0.250267  0.154959 -1.615 0.106299
WindGustSpeed 0.019025  0.031520  0.604 0.546122
WindSpeed9am  0.075650  0.049791  1.519 0.128669
WindSpeed3pm  0.007915  0.040858  0.194 0.846398
Humidity9am   0.054637  0.036961  1.478 0.139344
Humidity3pm   0.065631  0.038055  1.725 0.084594 .
Pressure9am   0.686808  0.235642  2.915 0.003561 ***
Pressure3pm  -0.824781  0.242893 -3.396 0.000685 ***
Cloud9am      0.044568  0.134685  0.331 0.740717
Cloud3pm      0.154399  0.143295  1.077 0.281261
Temp9am       0.126304  0.235915  0.535 0.592388
Temp3pm       0.229848  0.281252  0.817 0.413796
RainTodayYes  0.737751  0.890194  0.829 0.407244
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)
```

```

RainTodayYes 0.737751 0.890194 0.829 0.40/244
---
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 249.54 on 261 degrees of freedom
Residual deviance: 125.18 on 244 degrees of freedom
AIC: 161.18

Number of Fisher Scoring iterations: 7

> test_set<-test_set[-c(1,2,8,10,11,23)]
> mytestset2<-test_set[-18]
> test_set<-lapply(test_set,as.numeric)
Warning message:
In lapply(test_set, as.numeric) : NAs introduced by coercion
> predi1<-predict(model1,mytestset2,type="response")
> y_pred = ifelse(predi1 > 0.5, 1, 0)
> tab1 <- table(Predicted = predi1, Actual = test_set$RainTomorrow)
> missing_classerr <- mean(y_pred != test_set$RainTomorrow)
> print(paste('Accuracy =', 1 - missing_classerr))
[1] "Accuracy = 0.863636363636364"
> exp(coef(model1))
      (Intercept)      MinTemp      MaxTemp      Rainfall      Evaporation      Sunshine      WindGustSpeed
4.898238e+53  8.129726e-01  9.716381e-01  9.627377e-01  1.161268e+00  7.785930e-01  1.019207e+00
WindSpeed9am  WindSpeed3pm  Humidity9am  Humidity3pm  Pressure9am  Pressure3pm  Cloud9am
1.078585e+00  1.007946e+00  1.056157e+00  1.067833e+00  1.987361e+00  4.383310e-01  1.045576e+00
      Cloud3pm      Temp9am      Temp3pm      RainTodayYes
1.166957e+00  1.134627e+00  1.258409e+00  2.091226e+00
> anova(model1,test="Chisq")
Analysis of Deviance Table

Model: binomial, link: logit

Response: training_set$RainTomorrow

Terms added sequentially (first to last)


```

	Df	Deviance	Resid. Df	Resid. Dev	Pr(>Chi)
NULL			261	249.54	
MinTemp	1	12.179	260	237.36	0.0004834 ***
MaxTemp	1	8.516	259	228.84	0.0035200 **
Rainfall	1	0.382	258	228.46	0.5366904
Evaporation	1	0.384	257	228.08	0.5352358
Sunshine	1	32.168	256	195.91	1.414e-08 ***
WindGustSpeed	1	23.617	255	172.29	1.175e-06 ***
WindSpeed9am	1	0.015	254	172.28	0.9039217
WindSpeed3pm	1	0.020	253	172.26	0.8868929
Humidity9am	1	14.981	252	157.28	0.0001086 ***
Humidity3pm	1	4.125	251	153.15	0.0422614 *
Pressure9am	1	5.964	250	147.19	0.0145968 *
Pressure3pm	1	18.683	249	128.50	1.544e-05 ***
Cloud9am	1	0.035	248	128.47	0.8523197
Cloud3pm	1	1.240	247	127.23	0.2655565
Temp9am	1	0.561	246	126.67	0.4536685
Temp3pm	1	0.815	245	125.86	0.3667633
RainToday	1	0.676	244	125.18	0.4109843

```

---
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
> model2<-glm(training_set$RainTomorrow~MinTemp+MaxTemp+Sunshine+WindGustSpeed+Humidity9am+Humidity3pm+Pressure3pm+

```

```

> model2<-glm(training_set$RainTomorrow~MinTemp+MaxTemp+Sunshine+WindGustSpeed+Humidity9am+Humidity3pm+Pressure3pm,
())
> summary(model2)

Call:
glm(formula = training_set$RainTomorrow ~ MinTemp + MaxTemp +
    Sunshine + WindGustSpeed + Humidity9am + Humidity3pm + Pressure3pm,
    family = binomial(), data = mytrainset2)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-2.08109  -0.44194  -0.22296  -0.09943   2.44321

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept)  148.23544   50.24751   2.950 0.003177 **
MinTemp      -0.18863    0.07654  -2.464 0.013721 *
MaxTemp       0.28686    0.08491   3.379 0.000729 ***
Sunshine     -0.41413    0.11352  -3.648 0.000264 ***
WindGustSpeed  0.04813    0.02109   2.282 0.022474 *
Humidity9am   0.04856    0.02212   2.195 0.028157 *
Humidity3pm   0.02976    0.02473   1.203 0.228817
Pressure3pm  -0.15605    0.04862  -3.210 0.001328 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 249.54  on 261  degrees of freedom
Residual deviance: 142.88  on 254  degrees of freedom
AIC: 158.88

Number of Fisher Scoring iterations: 6

> 1-pchisq(deviance(model2)-deviance(model1),df.residual(model2)-df.residual(model1))
[1] 0.06023695
> predi2<-predict(model2,mytestset2,type="response")
> binca<-ifelse(binca=="Yes",c(1),c(0))
> summ=0
> for(i in 1:length(predi1)){summ=summ+(binca[i]-predi1[i])^2}
> summ
      4
20.91554
> summ=0
> for(i in 1:length(predi2)){summ=summ+(binca[i]-predi2[i])^2}
> summ
      4
20.19355
> mean(binca)
[1] 0.1829268
> sum=0
> for(i in 1:length(binca)){sum=sum+(binca[i]-mean(binca))^2}
> sum
[1] 49.02439
>

```